

ChiralFold: systematic detection of D-amino acid stereochemistry errors in the Protein Data Bank

Tommaso R. Marena^{1,*}

¹Department of Chemistry, The Catholic University of America, 620 Michigan Ave NE, 20064, Washington, DC, USA

*Corresponding author. marena@cua.edu

Abstract

Motivation: D-amino acids are essential in therapeutics and natural products, yet standard validators do not audit D-residue stereochemistry in deposited structures, and diffusion predictors violate D-chirality at high rates.

Results: We audited 12,573 D-residues in 4,623 PDB files and found 29 D-label/L-coordinate mismatches (0.23%) in 16 structures not covered by MolProbity's L-residue C α check. ChiralFold, calibrated on 362 structures (Ramachandran Spearman $\rho = 0.52$ vs wwPDB), detects errors via signed C α volume. Machine-checked Lean 4 proofs (Supplementary Notes S1) show that chirality sign cannot factor through pairwise distances alone.

Availability: <https://github.com/Tommaso-R-Marena/ChiralFold>; `pip install chiralFold (v3.5.1)`; **Python 3.9+**; **Linux/macOS/Windows**; **web UI** at <https://huggingface.co/spaces/The-Philosopher/ChiralFold-App> (Light/Dark; PDB/mmCIF/FASTA; optional RCSB and AlphaFold DB fetch).

Contact: marena@cua.edu

Supplementary information: Supplementary data, Lean 4 proofs with derivations, and extended analyses are available at Bioinformatics online and at <https://github.com/Tommaso-R-Marena/ChiralFold>.

Keywords D-amino acid, chirality, Protein Data Bank, structure validation, AlphaFold 3, Lean 4, formal verification

1. Introduction

D-amino acids resist proteolysis and underpin mirror-image therapeutics, nonribosomal peptide natural products, and bacterial cell-wall chemistry (Cava et al., 2011; Miyamoto and Homma, 2021). Mirror-image phage display has produced D-peptide binders with nanomolar affinity, including dPMI- γ ($K_d = 53$ nM against MDM2; PDB 3IWY) (Liu et al., 2010) and D-peptide HIV-entry inhibitors with crystallographically characterized complexes (Welch et al., 2009). Broader design strategies for D-peptide libraries and anti-amyloid constructs continue to expand the deposited D-ligand and D-peptide structural corpus (Funke and Willbold, 2010; Garton et al., 2018). As that corpus grows in the Protein Data Bank (PDB) (Berman et al., 2000; wwPDB consortium, 2019), the community needs validators that treat D-stereochemistry as a first-class quality criterion rather than an exotic edge case.

Two gaps are especially consequential. First, AlphaFold 3 (AF3) (Abramson et al., 2024)—building on the AlphaFold lineage (Jumper et al., 2021) and paralleling other deep structure predictors such as ESMFold and RoseTTAFold (Lin et al., 2023; Baek et al., 2021)—produces a 51% per-residue chirality violation rate on D-peptides, equivalent to random assignment (Childs et al., 2025), because its diffusion architecture denoises atom coordinates without hard stereochemical constraints. Second, MolProbity, the standard PDB quality tool (Chen et al., 2010; Lovell et al., 2003; Williams et al., 2018), validates C α chirality for standard L-amino acids but does not extend

that check to D-labeled residues. Annotation errors at D-sites can therefore persist undetected for decades, including through the 2006–2008 wwPDB remediation (Henrick et al., 2008). More generally, crystallographic validation has matured around geometry, density fit, and ligand chemistry (Read et al., 2011; Gore et al., 2017; Berman et al., 2016), yet D-label/coordinate consistency remains outside routine deposition checks.

We present ChiralFold, a pip-installable stereochemistry toolkit, and the first systematic PDB-wide audit of D-residue chirality. Using an independent signed-volume classifier grounded in CIP/IUPAC stereochemical convention (Favre and Powell, 2013), we scanned 12,573 checkable D-residues across 4,623 files (>91% coverage of each of the 18 standard D-amino acid Chemical Component Dictionary (CCD) codes (Shao et al., 2015)), classified error modes (genuine stereochemistry errors, CCD-code misassignments, polymer mislabels, and borderline near-zero volumes), and calibrated the auditor against wwPDB/MolProbity on an era-representative set of 362 structures. Critically, we complement the empirical survey with machine-checked Lean 4 no-go theorems showing that any geometric readout whose entire dependence factors through the pairwise distance matrix is necessarily blind to enantiomeric sign (Supplementary Notes S1; `formal/chirality_nogo/`). That formal obstruction motivates signed volume as the minimal non-distance stereochemical audit and places AF3-style distance-respecting denoising in a precise theoretical context—without claiming impossibility for general SE(3)-equivariant predictors that use angular or higher-order features.

Beyond auditing, ChiralFold provides exact L↔D mirroring, an AF3 chirality-correction post-processor, diastereomer enumeration, interface scoring, and optional online structure fetch (RCSB PDB IDs; AlphaFold DB by UniProt/AFDB identifier), released as open software (v3.5.1) with Colab notebooks, Docker images, a Hugging Face web UI (Light/Dark), and a five-minute offline reproduction path for the 29-error survey.

2. Materials and methods

2.1. Signed tetrahedron volume

Chirality at C_α is the signed volume of the tetrahedron formed by backbone atoms N, C_α , C, and C_β :

$$V = (\mathbf{N} - \mathbf{C}_\alpha) \cdot [(\mathbf{C} - \mathbf{C}_\alpha) \times (\mathbf{C}_\beta - \mathbf{C}_\alpha)]. \quad (1)$$

For correctly deposited D-residues, $V < 0$ by CIP convention (Favre and Powell, 2013); $V > 0$ at a D-label indicates L-coordinates. The classifier requires only four atom positions and uses no force field. Bond-length reference geometry for complementary auditor criteria follows Engh & Huber ideals (Engh and Huber, 1991).

The primary PDB survey was performed with an independent numpy-only script that imports no ChiralFold code (`benchmarks/independent_d_residue_verification.py`). Legacy PDB format was used for coverage statistics; mmCIF-only deposits were subsequently surveyed in a native-mmCIF expansion (Supplementary Methods): all 16 known-error structures re-verified (29/29 mismatches recovered), and the live mmCIF-only D-amino-acid universe ($n = 79$ at freeze) yielded one additional clear DLE mismatch (9BC4, $V = +2.82 \text{ \AA}^3$).

2.2. ChiralFold auditor and pipelines

The auditor validates six criteria: C_α chirality (Eq. 1), bond geometry (RMSD to Engh & Huber ideals), Ramachandran backbone dihedrals (hybrid empirical PDB grid plus calibrated rectangular regions), peptide-bond planarity (ω deviation from 180°), steric clashes (MolProbity-inspired overlap score with topology-aware exclusions; Supplementary Methods), and a weighted composite score (0–100). This design is intentionally complementary to MolProbity's all-atom L-protein validation suite (Chen et al., 2010; Williams et al., 2018): ChiralFold does not attempt to replace rotamer contour libraries or H-bond networks, but adds native D-residue chirality auditing that those tools omit. Clashscore is a secondary quality metric; the D-residue survey and Ramachandran calibration ($n = 362$) do not depend on it.

Mirror L↔D reflection $(x, y, z) \rightarrow (-x, y, z)$ inverts all stereocenters ($\det = -1$) while preserving bond lengths, angles, and torsion magnitudes exactly; residue names map to D-CCD equivalents (ALA→DAL, ...). Because reflection is an isometry, MolProbity-style clashscore is unchanged (`tests/test_clash_preservation.py`).

For AF3 or AlphaFold DB outputs, each C_α whose signed volume disagrees with its residue label is reflected across the plane of its three backbone neighbours (N, C, C_β), preserving CA–N/C/ C_β bond lengths to $< 10^{-3} \text{ \AA}$ and driving residual chirality violations to zero on synthetic fixtures (`ch`

`iralfold/af3_correct.py`; `tests/test_af3_correction.py`). Diastereomer enumeration builds all 2^k L/D patterns for k variable sites; interface scoring reports BSA, H-bonds, salt bridges, and a composite contact score. Optional RCSB/AlphaFold DB fetch and PDB/mmCIF/FASTA inputs are documented in Supplementary Methods.

2.3. Ramachandran calibration against wwPDB

We assembled an era-representative stratified sample spanning X-ray ultra ($\leq 1.20 \text{ \AA}$), high ($1.20\text{--}1.80 \text{ \AA}$), medium ($1.80\text{--}2.50 \text{ \AA}$), and low ($2.50\text{--}3.40 \text{ \AA}$) resolution tiers, solution NMR, and cryo-EM, with deposition dates drawn uniformly within each bin (seed 20260612). Structures were obtained in legacy PDB format where available; mmCIF-only entries were converted with chain-ID remapping (A–Z, then 0–9) that prevents multi-character chain-ID truncation. Entries with >36 distinct chains were skipped. For each structure, Ramachandran outlier percentages from wwPDB validation XML were compared to ChiralFold auditor output, following the spirit of method-agreement analysis used elsewhere in measurement science (Bland and Altman, 1986). One glycopeptide antibiotic (5M2K, vancomycin– Zn^{2+}) was excluded under a pre-specified criterion (`results/5m2k_benchmark_exclusion.json`; `wwpdb_rama_outlier_pct > 50%`), yielding $n = 362$ of 363 successfully audited entries. Script: `benchmarks/expand_ramachandran_benchmark.py`.

2.4. Formal verification in Lean 4

To show that signed volume is not merely empirically useful but formally necessary for distinguishing enantiomers under a distance-only information constraint, we formalized a family of no-go theorems in Lean 4 (v4.28.0) with Mathlib (`formal_chirality_nogo/`). A four-point configuration $P \in (\mathbb{R}^3)^4$ has pairwise squared distance matrix `distMatrix4(P)` and oriented volume

$$\text{orient4}(P) = \det[\mathbf{b} - \mathbf{a} \mid \mathbf{c} - \mathbf{a} \mid \mathbf{d} - \mathbf{a}]. \quad (2)$$

Reflection through the xy -plane preserves every entry of `distMatrix4` but negates `orient4`. A model M *factors through distances* if $M(P) = F(\text{distMatrix4}(P))$ for some F ; such models are provably reflection-invariant and therefore cannot recover chirality sign on all configurations. The Gram identity $\text{orient4}^2 = \det(G)$ shows that distances determine magnitude but not sign—precisely the information Eq. 1 supplies at C_α .

All theorems compile without `sorry` or custom axioms beyond Lean's standard library (lake build). Full informal derivations, theorem-by-theorem biological interpretation, concrete tetrahedron examples, and the SE(3) scope limitation appear in Supplementary Notes S1.

Informally: (1) reflection is an isometry, so it preserves `distMatrix4`; (2) $\det = -1$, so it negates `orient4`; (3) any $M = F \circ \text{distMatrix4}$ is therefore reflection-invariant; (4) the standard tetrahedron and its mirror share distances but opposite orientation, so no such M can recover signed chirality on all inputs; (5) the Gram identity shows distances recover magnitude only. ChiralFold's Eq. 1 supplies the missing sign at C_α .

2.5. Software, availability, and reproducibility

ChiralFold v3.5.1 is installable via `pip install chiralFold`, `conda (conda-recipe/)`, `Docker (ghcr.io/tommaso-r-marena/chiralFold:3.5.1)`, and a Gradio web UI (local `chiralFold-web` or Hugging Face Space with Light/Dark theme, PDB/mmCIF/FASTA upload, and optional RCSB/AlphaFold DB fetch). Fetch is download-only: ChiralFold does not run AF2, Boltz, or AF3 inference. Reviewers can reproduce the 29-error survey offline after clone with `python benchmarks/reproduce_d_residue_errors.py` (numpy only). Lean proofs build with `lake build` under `formal/chirality_nogo/`. Colab notebooks cover the five-minute survey, mmCIF expansion, experimental validation, unusual cases, and the interactive results dashboard.

3. Results

3.1. PDB D-residue survey

We verified 12,573 D-amino acid residues with complete backbone atoms (N, C_α , C, C_β) across 4,623 PDB files, covering >91% of all RCSB entries for each of the 18 standard D-amino acid CCD codes (Shao et al., 2015). Of these, 12,544 (99.77%) showed correct D-chirality ($V < 0$) and 29 (0.23%) showed L-chirality ($V > 0$) despite D-labels (Fig. 1a). An independent numpy-only script (no ChiralFold imports) produced the same classifications.

The 29 mismatches occur in 16 PDB structures and fall into four classes (Supplementary Table S1): **(i) Stereochem** (6 errors / 6 structures)—biology requires D, coordinates show L (e.g. 1ABI hirulog DPN, 1HHZ cell-wall DAL at 0.99 Å, 1XT7 daptomycin DSG, 2RMI astressin DPN); **(ii) CCD-Code** (18 errors / 5 structures)—coordinates are L but the depositor selected a D-form CCD code (e.g. 1OF6 eight DTY←L-Tyr; 2ATS three DLY despite a title stating “(S)-lysine”; 3RIT five DLY); **(iii) Mismatch** (3 errors)—D-codes on polymer residues in L-proteins (1MCB, 1UHG, 2AOU); **(iv) Borderline** (2 errors)—near-zero volumes ($|V| < 1 \text{ Å}^3$; 1KO0, 2H9E) retained for transparency but flagged separately.

An automated validation pipeline (`benchmarks/experimental_structure_validation.py`) cross-checks all 16 structures against RCSB experimental metadata and wwPDB CCD InChI stereochemistry: 14/14 non-borderline cases pass. Errors concentrate in nine CCD codes; nine codes are confirmed clean at zero errors (Fisher’s exact $p = 1.96 \times 10^{-10}$; $\chi^2 = 113.78$, $p = 2.33 \times 10^{-16}$ across 18 codes). DTY and DLY alone account for 19 of 29 mismatches (Fig. 1a). MolProbity’s standard C_α chirality check applies to L-residues only and does not flag these D-label mismatches (Chen et al., 2010; Lovell et al., 2003). Thirteen of 16 error structures predate 2006 wwPDB remediation (Mann–Whitney $p = 1.14 \times 10^{-4}$; Fig. 1b) (Henrick et al., 2008); resolution does not predict errors at the structure level ($p = 0.19$).

Signed-volume classification recovers errors in cyclic peptides (1XT7, 2RMI), non-standard CCD ligands (1OF6), and ultra-high-resolution deposits (1HHZ), showing the method is not limited to canonical globular proteins or low-density maps

Table 1. Selected machine-checked no-go theorems (Lean 4; all sorry-free). Full names and informal proofs: Supplementary Notes S1.

Theorem (short name)	Geometric / biological consequence
<code>same_distances_opposite_orientations_standardTet</code>	Enantiomeric tetrahedra can share all distances yet opposite signed volume.
<code>det_gram3_eq_orient_sq</code>	Distances fix $ \text{orient} ^2$ via the Gram determinant, not the sign.
<code>improper_orthogonal_distance_chirality_obstruction</code>	Improper orthogonal maps preserve distances and reverse orientation.
<code>chiralitySignBool_not_distance_factorable</code>	No distance-matrix-only classifier recovers chirality sign universally.
<code>proteinCaOrient_not_function_of_six_pairwise_distances</code>	(N, C_α , C, C_β) signed volume needs information beyond six distances.

(Supplementary Methods). The native-mmCIF expansion re-verified all 16 known-error structures (29/29) and scanned the live mmCIF-only D-AA universe ($n = 79$), finding one additional clear DLE mismatch (9BC4). Together these results place the historical D-label problem on a quantitative footing: rare at the residue level (0.23%), non-random by CCD code, enriched in pre-remediation deposits, and invisible to L-centric validation pipelines still used at deposition time (Gore et al., 2017; Read et al., 2011).

3.2. MolProbity calibration

On $n = 362$ structures, ChiralFold Ramachandran outlier percentages agree with wwPDB (Spearman $\rho = 0.52$, 95% CI [0.42, 0.61], $p = 1.5 \times 10^{-26}$; mean 0.78% vs 0.83%; Fig. 2). Pearson r is outlier-sensitive (0.33 for $n = 363$; 0.85 after excluding 5M2K); Spearman ρ is the primary agreement metric. A Bland–Altman view of the same comparison is provided in Supplementary Fig. S3 (Bland and Altman, 1986). MolProbity remains stronger in Ramachandran contour precision and rotamer completeness (Williams et al., 2018); ChiralFold adds native D-residue chirality auditing that MolProbity does not provide. ChiralFold rotamer scoring is currently chi1-only.

3.3. Lean 4 no-go theorems and AF3

The Lean formalization proves several layers of obstruction (Table 1; full catalog in Supplementary Notes S1). The central witness is the standard tetrahedron with vertices at 0, e_1 , e_2 , e_3 : `orient4` = +1, while its z -reflection has `orient4` = −1 with an identical distance matrix. Consequently no Boolean classifier on distances alone can separate the pair, and no continuous representation that factors through `distMatrix4` can recover chirality sign on all configurations. Specializing atom order to (N, C_α , C, C_β) yields the protein-facing statement that C_α signed volume is not a function of those six distances.

Scope.

These theorems characterize *distance-only* representations. They are not an impossibility result for AF3 (Abramson et al., 2024) or for general SE(3)-equivariant networks that may use angular, tensorial, or sequence-derived features (Supplementary Notes S1). ChiralFold does not model thermodynamic ensembles or molecular-mechanics free energies; it enforces an algebraic signed-volume constraint during coordinate mapping and auditing.

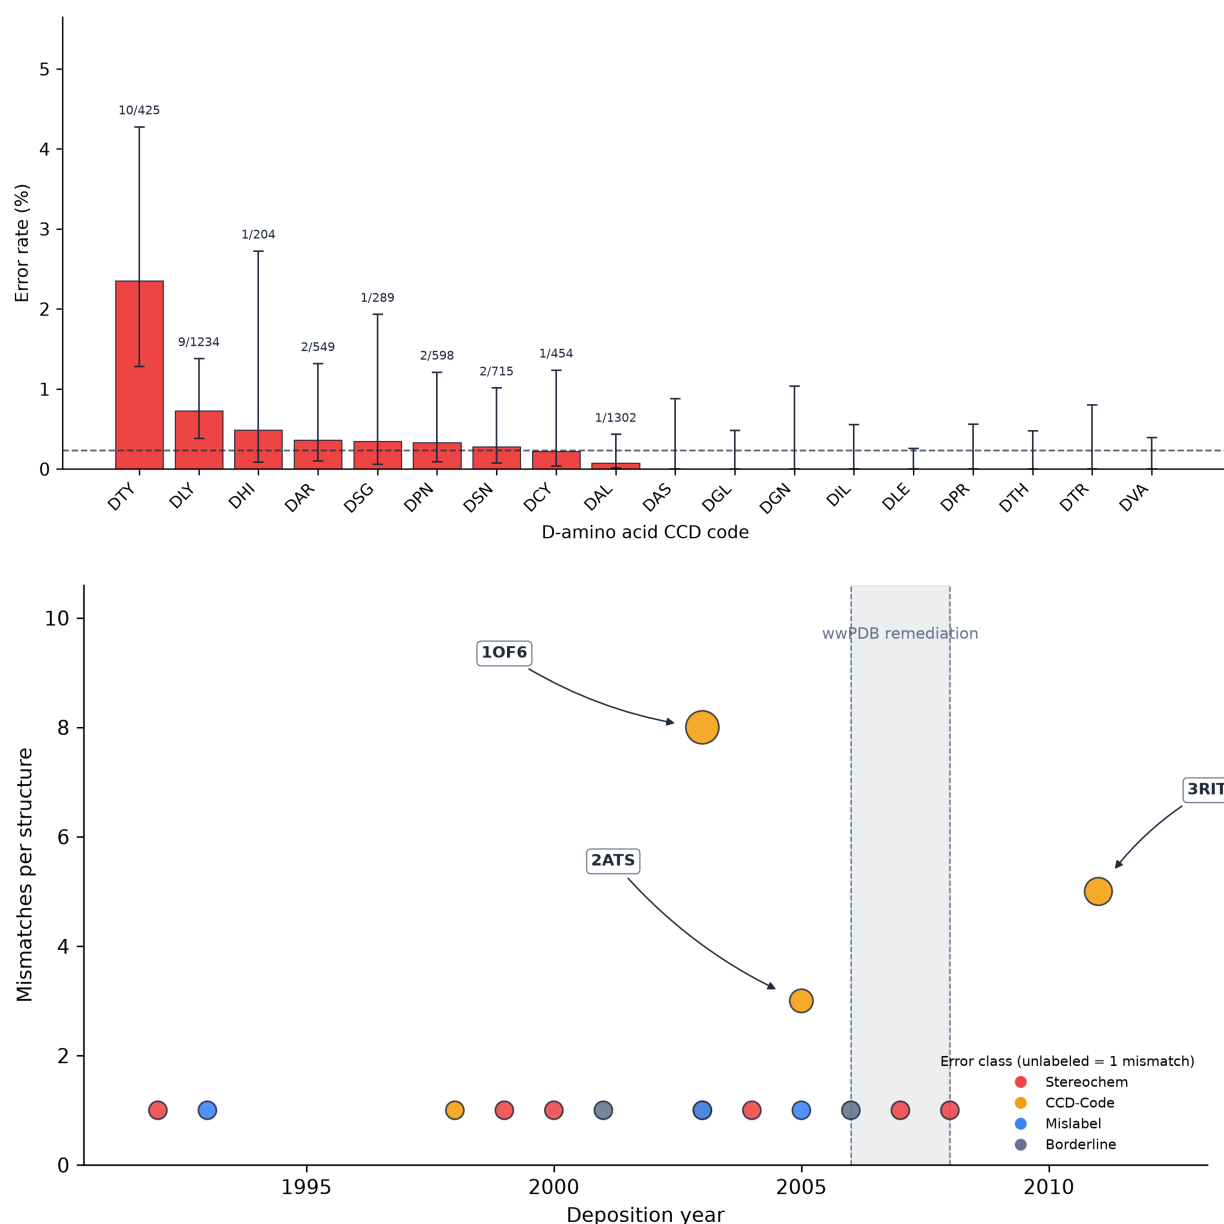


Figure 1 (a) Per-CCD mismatch rate with Wilson 95% CIs; red: codes with ≥ 1 error (labels show errors/residues checked); teal: confirmed clean (unlabeled; coverage in SI). Dashed line: dataset-wide mean (0.23%). (b) Mismatches per error structure vs deposition year (marker size $\propto$ count; color = error class). Shaded band: 2006–2008 wwPDB remediation. Multiplicity outliers (10F6, 2ATS, 3RIT) are CCD-code misassignments; single-mismatch structures are unlabeled for clarity.

3.4. AF3 D-peptide sequences and synthetic correction

Using D-peptide sequences from Childs et al. (2025)—DEHELLETAARWFYEIAKR (DP19), LWQHEATWK (DP9), and DWWPLAFEALLR (DP12)—ChiralFold produces 0/478 chirality violations across 41 sequence/pattern combinations. This 0% rate is guaranteed by construction: each residue is built with explicit SMILES stereochemistry. AF3’s 51% rate reflects diffusion treating D-residues as noise. Comparing ChiralFold’s 467/0 correct/violated residues on these sequences to Childs et al.’s reported AF3 universe counts yields Fisher’s exact $p \approx 6.7 \times 10^{-144}$. This is a construction-versus-learning contrast,

not a head-to-head fold-accuracy comparison: ChiralFold does not predict folds, and is complementary to AF3, ESMFold, and RoseTTAFold as a stereochemistry post-processor (Abramson et al., 2024; Lin et al., 2023; Baek et al., 2021).

A synthetic AF3-mimetic resource benchmark further shows 100% detection/correction recall, 0% residual violations, and a 0.0367 s local correction pipeline across three Childs-system peptides (Supplementary Fig. S4), supporting correction as lightweight stereochemistry post-processing rather than AF3 re-inference. Mirror-image transformation of the p53:MDM2 complex (1YCR) preserves contact geometry to 0.001 Å by construction; experimental K_d measurement is required to

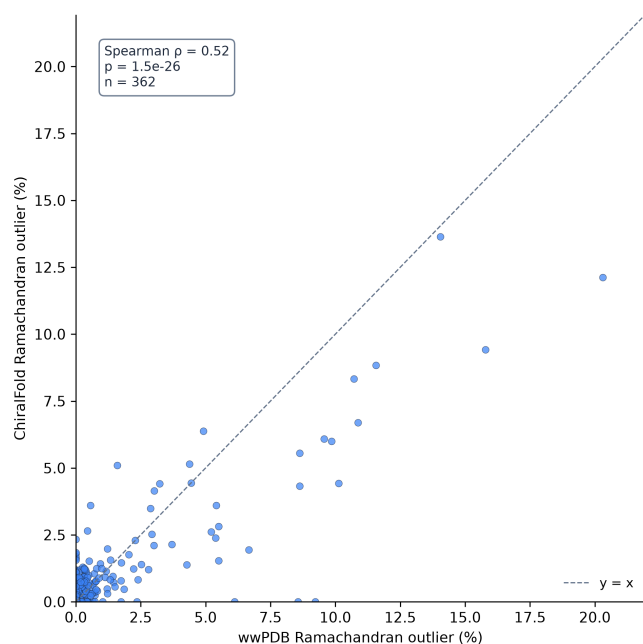


Figure 2 ChiralFold vs wwPDB Ramachandran outlier percentages on the era-representative benchmark ($n = 362$ after excluding 5M2K). Dashed line: $y = x$. Spearman $\rho = 0.52$ ($p = 1.5 \times 10^{-26}$).

validate binding and is outside this computational study (Liu et al., 2010).

4. Discussion

The signed $C\alpha$ volume distinguishes three operationally distinct annotation failures without biological context. Genuine stereochemistry errors (1ABI, 1D7T, 1HHZ, 1XT7, 2RMI, 2W76) are the most consequential: biology constrains the residue to D, yet deposited coordinates show L with $|V|$ typically 1.85–2.77 Å³. The 1HHZ case is particularly telling—at 0.99 Å atomic resolution, D-Ala in a cell-wall pentapeptide carries $V = +2.70$, the largest positive value in the dataset, in a context where D-alanine is a biosynthetic requirement of peptidoglycan cross-linking (Cava et al., 2011; Miyamoto and Homma, 2021). CCD-code misassignments (1BG0, 1OF6, 1P52, 2ATS, 3RIT) define a second mode: coordinates correctly encode L-stereochemistry, but the depositor selected the D-form CCD code. Because coordinates are correct, remediation is a metadata fix, not re-refinement. Polymer mislabels (1MCB, 1UHG, 2AOU) apply D-codes to residues in L-proteins, sometimes despite COMPND records specifying L-form.

The CCD category has a clear upstream fix. The wwPDB Chemical Component Dictionary encodes configuration in the InChI stereo layer (/m0 = S/L; /m1 = R/D) (Shao et al., 2015; Westbrook et al., 2015). An automated deposition check that computes Eq. 1 and compares its sign to the selected CCD /m flag would have caught every CCD-code case reported here before archival acceptance (Berman et al., 2016; Gore et al., 2017). The 2ATS title field literally states “(S)-lysine” while depositing DLY—a pipeline gap, not a judgment call. Such a check is inexpensive relative to existing ligand-validation

workshops and wwPDB OneDep geometry screens (Berman et al., 2016; Read et al., 2011).

Deposition year predicts errors (logistic OR= 1.60 per year; Mann–Whitney $p = 1.14 \times 10^{-4}$), whereas resolution does not ($p = 0.19$), consistent with labeling/convention failures from an era without D-specific validation rather than density-quality artifacts (Henrick et al., 2008). Errors concentrate in DTY/DLY/DPN/DSN/DAR while nine codes remain clean at >91% coverage—a non-random clustering (Fisher $p = 1.96 \times 10^{-10}$) that supplies a practical checklist for depositors and curators.

The Lean no-go results explain why distance-respecting validation and diffusion post-processing cannot substitute for signed-volume audits: pairwise distances determine $|\text{orient}|^2$ but not its sign (Supplementary Notes S1). ChiralFold's Eq. 1 is the minimal geometric invariant that closes this gap for $C\alpha$. ChiralFold is complementary to MolProbity, not a replacement: MolProbity remains stronger for L-protein Ramachandran contours and rotamers (Williams et al., 2018); ChiralFold addresses the D-label blind spot that left the 29 mismatches invisible to existing pipelines.

Therapeutically, correct D-stereochemistry matters for proteolysis resistance and target engagement in D-peptide programs (Welch et al., 2009; Funke and Willbold, 2010; Garton et al., 2018). Structurally, silent D-label errors corrupt any downstream analysis that trusts residue names as stereochemical ground truth—including training sets for structure predictors and automated CCD assignment tools. ChiralFold is a stereochemistry post-processor, not a de novo folder: use AF3, ESMFold, or RoseTTAFold for folding (Abramson et al., 2024; Lin et al., 2023; Baek et al., 2021), then audit or correct with ChiralFold.

Scope and limitations.

ChiralFold does **not** claim to model the complete thermodynamic landscape of structural ensembles, sample Boltzmann-weighted conformers, or replace molecular-mechanics force fields. Its operations fix an explicit *algebraic data-representation constraint* during coordinate mapping: signed $C\alpha$ tetrahedron volume for auditing, exact improper reflection for L↔D mirroring and AF3 correction, and deterministic diastereomer enumeration. Evaluating the tool as an energy-minimization or ensemble-dynamics simulator would misstate its design. Additional limitations include chi1-only rotamer scoring, two borderline cases with $|V| < 1.0$ Å³, no experimental binding validation of mirror-designed peptides, a MolProbity-inspired clashscore that places Reduce-style HN/HA hydrogens and applies geometry-gated H-bond suppression but is not bit-identical to Probe (Supplementary Methods), and the primary survey's use of legacy PDB format (closed via mmCIF re-verification of known errors and the live mmCIF-only universe survey; Supplementary Methods). Formal theorems characterize distance-only representations and do not claim impossibility for general SE(3)-equivariant predictors.

5. Conclusion

ChiralFold combines a PDB stereochemistry auditor calibrated against wwPDB/MolProbity ($n = 362$; Spearman $\rho = 0.52$), machine-checked Lean 4 proofs that distance-only models

cannot determine chirality sign (Supplementary Notes S1), exact L↔D mirroring, and an AF3 chirality-correction pipeline. Applied to 12,573 D-residues, the auditor identifies 29 D-label/L-coordinate mismatches in 16 deposited structures—errors invisible to existing validation pipelines—and motivates a concrete deposition recommendation: validate every standard-amino-acid stereocenter against the configuration encoded in the selected CCD code at deposit time. The check requires one cross product, one dot product, and one sign comparison per residue. Code, data, formal proofs, and reproduction scripts are publicly available (v3.5.1).

Data availability

Verification data: <https://github.com/Tommaso-R-Marena/ChiralFold> (results/d_residue_verification.csv). Experimental validation: results/experimental_validation_report.json. mmCIF survey outputs: results/mmCIF_d_residue_expansion_*. Lean proofs: formal/chirality_nogo/. Ramachandran benchmark: results/ramachandran_279struct_chainfix_*. MolProbity head-to-head panel (secondary clash/planarity fields regenerated after topology-aware clashscore): results/molprobity_comparison.json. Software: PyPI chiralFold (v3.5.1); Docker image chiralFold:3.5.1 on GHCR; GitHub release v3.5.1; Hugging Face Space <https://huggingface.co/spaces/The-Philosopher/ChiralFold-App>.

Funding

None declared.

Conflict of interest

None declared.

Author contributions

T.R.M. conceived the study, implemented ChiralFold, performed all analyses, formalized proofs, and wrote the manuscript.

References

- J. Abramson et al. Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*, 630:493–500, 2024. doi: 10.1038/s41586-024-07487-w.
- M. Baek et al. Accurate prediction of protein structures and interactions using a three-track neural network. *Science*, 373: 871–876, 2021. doi: 10.1126/science.abj8754.
- H. Berman et al. Outcome of the first wwPDB/ccdc/d3r ligand validation workshop. *Structure*, 24:502–508, 2016. doi: 10.1016/j.str.2016.02.017.
- H. M. Berman et al. The Protein Data Bank. *Nucleic Acids Res.*, 28:235–242, 2000. doi: 10.1093/nar/28.1.235.
- J. M. Bland and D. G. Altman. Statistical methods for assessing agreement between two methods of clinical measurement. *Lancet*, 327:307–310, 1986. doi: 10.1016/S0140-6736(86)90837-8.
- F. Cava et al. Distinct pathways for modification of the bacterial cell wall by non-canonical D-amino acids. *EMBO J.*, 30: 3442–3453, 2011. doi: 10.1038/emboj.2011.246.
- V. B. Chen et al. MolProbity: all-atom structure validation for macromolecular crystallography. *Acta Crystallogr. D*, 66: 12–21, 2010. doi: 10.1107/S0907444909042073.
- C. M. Childs, P. Zhou, and B. R. Donald. Has AlphaFold 3 solved the protein folding problem for D-peptides? *bioRxiv*, 2025. doi: 10.1101/2025.03.14.643307.
- R. A. Engh and R. Huber. Accurate bond and angle parameters for X-ray protein structure refinement. *Acta Crystallogr. A*, 47:392–400, 1991. doi: 10.1107/S0108767391001071.
- H. A. Favre and W. H. Powell. Nomenclature of Organic Chemistry: IUPAC recommendations and preferred names 2013. *RSC Publishing*, 2013. Blue Book; CIP sequence rules.
- S. A. Funke and D. Willbold. Peptides designed to prevent A β aggregation and neurotoxicity. *Curr. Pharm. Des.*, 16: 1679–1690, 2010. doi: 10.2174/138161210791164153.
- M. Garton et al. Methods for generating and screening libraries of D-peptides. *Curr. Opin. Struct. Biol.*, 51:123–131, 2018. doi: 10.1016/j.sbi.2018.04.001.
- S. Gore et al. Validation of structures in the Protein Data Bank. *Structure*, 25:1916–1927, 2017. doi: 10.1016/j.str.2017.10.009.
- K. Henrick et al. Remediation of the protein data bank archive. *Nucleic Acids Res.*, 36:D426–D433, 2008. doi: 10.1093/nar/gkm937.
- J. Jumper et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596:583–589, 2021. doi: 10.1038/s41586-021-03819-2.
- Z. Lin et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379:1123–1130, 2023. doi: 10.1126/science.ade2574.
- M. Liu et al. D-peptide inhibitors of the p53–MDM2 interaction for targeted molecular therapy of malignant neoplasms. *Proc. Natl. Acad. Sci. U.S.A.*, 107:14321–14326, 2010. doi: 10.1073/pnas.1008930107.
- S. C. Lovell et al. Structure validation by C α geometry: ϕ , ψ and C β deviation. *Proteins*, 50:437–450, 2003. doi: 10.1002/prot.10286.
- T. Miyamoto and H. Homma. D-amino acid metabolism in bacteria. *J. Biochem.*, 170:5–13, 2021. doi: 10.1093/jb/mvab043.
- R. J. Read et al. A new generation of crystallographic validation tools for the protein data bank. *Structure*, 19:1395–1412, 2011. doi: 10.1016/j.str.2011.08.006.
- C. Shao et al. The chemical component dictionary: complete descriptions of constituent molecules in experimentally determined 3d macromolecules in the Protein Data Bank. *Bioinformatics*, 31:1274–1278, 2015. doi: 10.1093/bioinformatics/btu789.
- B. D. Welch et al. Structure of a D-peptide complex that inhibits HIV entry. *Proc. Natl. Acad. Sci. U.S.A.*, 106:13330–13335, 2009. doi: 10.1073/pnas.0903939106.
- J. D. Westbrook et al. The chemical component dictionary: complete descriptions of constituent molecules in experimentally determined 3D macromolecules in the Protein Data Bank. *Bioinformatics*, 31:1274–1278, 2015. Same CCD resource as Shao et al.; cited for deposition tooling context.

C. J. Williams et al. MolProbity: More and better reference data for improved all-atom structure validation. *Protein Sci.*, 27: 293–315, 2018. doi: 10.1002/pro.3330.

wwPDB consortium. Protein Data Bank: the single global archive for 3D macromolecular structure data. *Nucleic Acids Res.*, 47:D520–D528, 2019. doi: 10.1093/nar/gky949.